

```
%pip install -q otter-grader
```

```
import sys

IN_COLAB = 'google.colab' in sys.modules

if IN_COLAB:
    ! git clone https://github.com/tanish738/CDA-HW2-TESTS.git tests
    import otter
    grader = otter.Notebook()
else:
    print("Not running in Colab")
```

fatal: destination path 'tests' already exists and is not an empty directory.

```
import otter
grader = otter.Notebook()
```

Assignment Summary

In this assignment, we explored and implemented several **fundamental machine learning algorithms from scratch**. Each method approached classification differently, giving us insights into linear, tree-based, instance-based, and anomaly-detection models.

Distribution of Problems Solved

1. Logistic Regression

- Implemented sigmoid, cost function, gradient computation, parameter updates, and prediction.
- Learned how linear models classify data using probabilities and decision boundaries.

2. One-Class SVM

- Implemented anomaly detection for imbalanced datasets.
- Learned how to separate “normal” data from potential outliers using a margin-based approach.

3. Decision Trees

- Implemented Gini impurity and best-split search.
- Understood how recursive partitioning builds interpretable, rule-based models.

4. K-Nearest Neighbors (KNN)

- Implemented distance calculation, neighbor search, and majority voting.
- Learned how instance-based models adapt to complex, non-linear decision boundaries.

✓ Logistic Regression from Scratch (Binary Classification)

Assignment Overview

In this assignment, you will implement a **binary logistic regression classifier** entirely from scratch using **Python** and the **NumPy** library.

The goal is to **classify data points into one of two distinct classes** (0 or 1) by learning the underlying decision boundary.

This exercise will strengthen your understanding of the **mathematics and mechanics** behind one of the most fundamental machine learning algorithms.

Learning Objectives

By completing this assignment, you will:

- Understand how logistic regression works under the hood.
- Implement essential components step by step:

1. **Sigmoid function** – to map raw scores into probabilities ✦

2. **Cost function** – binary cross-entropy loss.
 3. **Gradient computation** – calculate how to adjust parameters.
 4. **Gradient Descent update rule** – iteratively optimize model parameters.
 5. **Prediction function** – classify samples based on learned weights.
- Train and evaluate your logistic regression classifier on a dataset.


```
import sys
import numpy as np
import matplotlib.pyplot as plt
!curl -so p0_data.txt "https://dl.dropboxusercontent.com/scl/fi/8rhvgae0al2s9z9oabxoo/p0_data.txt?rlkey=bj4nbz4q013vz0a
data = np.loadtxt('p0_data.txt', delimiter=',')

train_X = data[:, 0:2]
train_y = data[:, 2]

# Get the number of training examples and the number of features
m_samples, n_features = train_X.shape
print("# of training examples = ", m_samples)
print("# of features = ", n_features)

pos = np.where(train_y == 1)
neg = np.where(train_y == 0)
plt.scatter(train_X[pos, 0], train_X[pos, 1], marker='o', c='b')
plt.scatter(train_X[neg, 0], train_X[neg, 1], marker='x', c='r')
plt.xlabel('Exam 1 score')
plt.ylabel('Exam 2 score')
plt.legend(['Admitted', 'Not Admitted'])
plt.show()
```

```
# of training examples = 100
# of features = 2
```



```
def sigmoid(z: int)-> float:
    """
    Sigmoid function
    Parameters
    -----
    z : float or numpy.ndarray
        The input value(s) for which the sigmoid is to be calculated. The function is designed to handle both single nu

    Returns
    -----
    float or numpy.ndarray
        The result of the sigmoid calculation. The output will be of the same type as the input `z`, with all values ma
    """
    # ***** ENTER CODE *****
    s = 1 / (1 + np.exp(-z))
    return s
    # ***** END CODE *****
```

```
grader.check("q1")
```

q1 passed! 🌟

✓ Logistic Regression Cost Function

We will implement the **Binary Cross-Entropy Loss** (Log Loss) for logistic regression.

1. Use the **sigmoid function** for predictions:
2. **Clip predicted probabilities** to avoid $\log(0)$ errors:
3. **Compute the average loss** over all training examples:

```
def cost_function(theta: np.ndarray, X: np.ndarray, y: np.ndarray) -> float:
    """ Logistic regression cost function
    Hint: clip h to avoid log(0).
    Parameters
    -----
    theta : numpy.ndarray
        The model's parameters (weights and bias) with shape (n_features + 1,).

    X : numpy.ndarray
        The feature matrix with shape (m_samples, n_features + 1).

    y : numpy.ndarray
        The target vector containing the true labels (0s and 1s) with shape (m_samples,).

    Returns
    -----
    float
        The computed logistic regression cost, a single scalar value.
    """

    m = len(y)
    # ***** ENTER CODE *****
    h = sigmoid(np.dot(X, theta))
    # Clip predictions to avoid log(0)
    h = np.clip(h, 1e-15, 1 - 1e-15)
    cost = -(1/m) * (np.dot(y, np.log(h)) + np.dot((1-y), np.log(1-h)))
    return cost
    # ***** END CODE *****
```

```
grader.check("q2")
```

q2 passed! 🤖

✓ Gradient Update

To compute the gradient of the loss function with respect to the parameters:

1. **Subtract actual labels from predictions:**
2. **Multiply by the feature matrix transpose:**
3. **Divide by the number of examples (m):**

```
def gradient_update(theta: np.ndarray, X: np.ndarray, y: np.ndarray) -> np.ndarray:
    """ Gradient for logistic regression
    Parameters
    -----
    theta : numpy.ndarray
        The model's current parameters (weights and bias) with shape `(n_features + 1,)`.

    X : numpy.ndarray
        The feature matrix with shape `(m_samples, n_features + 1)`.

    y : numpy.ndarray
        The target vector containing the true labels (0s and 1s) with shape `(m_samples,)`.

    Returns
    -----
```

```

numpy.ndarray
    The computed gradient vector with the same shape as `theta`, `(n_features + 1)`.

"""
m = len(y)
# ***** ENTER CODE *****
h = sigmoid(np.dot(X, theta))
gradient = (1/m) * np.dot(X.T, (h - y))
return gradient
# ***** END CODE *****

```

```
grader.check("q3")
```

q3 passed! 🎉

✓ Gradient Descent Algorithm

In each iteration of Gradient Descent we do two key updates:

1. Gradient Update

- Compute the gradient of the loss function w.r.t. the parameters.

2. Parameter Update

- Update parameters in the opposite direction of the gradient:

```

def gradient_descent(theta, X, y, alpha, max_iterations, print_iterations):
    """ Batch gradient descent algorithm """
    iteration = 0
    prev_cost = float('inf') # Track previous cost

    while iteration < max_iterations:
        iteration += 1
        # ***** ENTER YOUR CODE *****
        gradient = gradient_update(theta, X, y)
        theta = theta - alpha * gradient
        # ***** END CODE *****

        # For every print_iterations number of iterations
        if iteration % print_iterations == 0 or iteration == 1:
            cost = cost_function(theta, X, y)
            print("[ Iteration", iteration, "]", "cost =", cost)

            # Visualization
            plt.figure(figsize=(5, 4))
            plt.xlim([20, 110])
            plt.ylim([20, 110])

            pos = np.where(y == 1)
            neg = np.where(y == 0)

            # Plot original data
            plt.scatter(X[pos, 1], X[pos, 2], marker='o', c='b')
            plt.scatter(X[neg, 1], X[neg, 2], marker='x', c='r')
            plt.xlabel('Exam 1 score')
            plt.ylabel('Exam 2 score')
            plt.legend(['Admitted', 'Not Admitted'])

            # Plot decision boundary
            t = np.linspace(20, 110, 100)
            if abs(theta[2]) > 1e-6:
                decision_boundary = -(theta[0] + theta[1] * t) / theta[2]
                mask = (decision_boundary >= 20) & (decision_boundary <= 110)
                if np.any(mask):
                    plt.plot(t[mask], decision_boundary[mask], c='g', linewidth=2, label='Decision Boundary')

            plt.title(f'Iteration {iteration}')
            plt.show()

        # Early stopping if cost stops improving significantly
        if iteration > 1000 and abs(prev_cost - cost) < 1e-6:
            print(f"Converged at iteration {iteration}")

```

```

        break

    prev_cost = cost

    return theta

```

```
grader.check("q4")
```

q4 passed! 🎉

✓ Predict Function (Logistic Regression)

Steps:

1. Compute probabilities using the **sigmoid function**:
2. Apply **threshold at 0.5**:
 - If $(h \geq 0.5) \rightarrow$ predict **1**
 - If $(h < 0.5) \rightarrow$ predict **0**

✓ Theory Question 1 - Accuracy of Logistic Regression

Feel Free to change the parameters : [Hint - Use the visualization to get insights of data]

- The accuracy expected is 90 and you might have to change the parameters to get this
- This will be manually graded we will also check the fit of the line in the graphs
- In the submission for this question you need to present all the graphs basically the output of the cell below after it is executed

```

train_X_with_bias = np.column_stack([np.ones(len(train_X)), train_X])
print(f"Data shape with bias: {train_X_with_bias.shape}")

initial_theta = np.array([0, 0, 0]) # you can change the initialization for better results
alpha_test = 0.01
max_iter = 1000
print_iter = 100

learned_theta = gradient_descent(initial_theta, train_X_with_bias, train_y, alpha_test, max_iter, print_iter)

def predict(theta, X):
    pass
    # Predict labels using learned parameters
    # ***** ENTER CODE *****
    probabilities = sigmoid(np.dot(X, theta))
    predictions = (probabilities >= 0.5).astype(int)
    return predictions
    # ***** END CODE *****

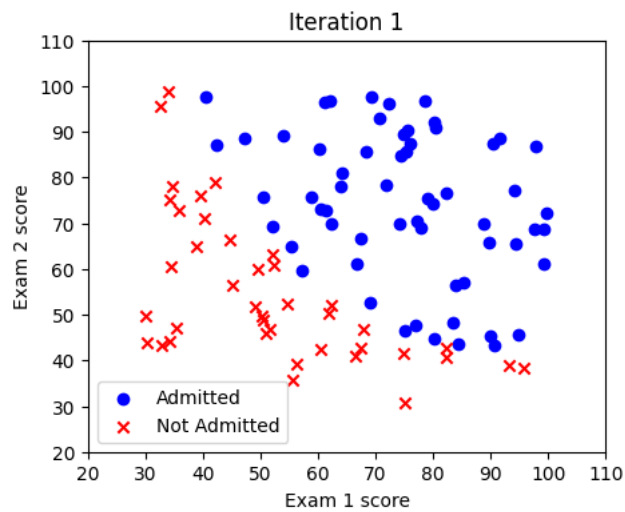
def calculate_accuracy(theta, X, y):
    """ Calculate classification accuracy """
    predictions = predict(theta, X)
    accuracy = np.mean(predictions == y) * 100
    return accuracy

training_accuracy = calculate_accuracy(learned_theta, train_X_with_bias, train_y)
print(f"\nTraining Accuracy: {training_accuracy}%")


```




```
Data shape with bias: (100, 3)
[ Iteration 1 ] cost = 4.9606001433589375
```




```
[ Iteration 100 ] cost = 5.264256705792918
```




```
[ Iteration 200 ] cost = 7.867373833565725
```




```
[ Iteration 300 ] cost = 12.484114448455761
```

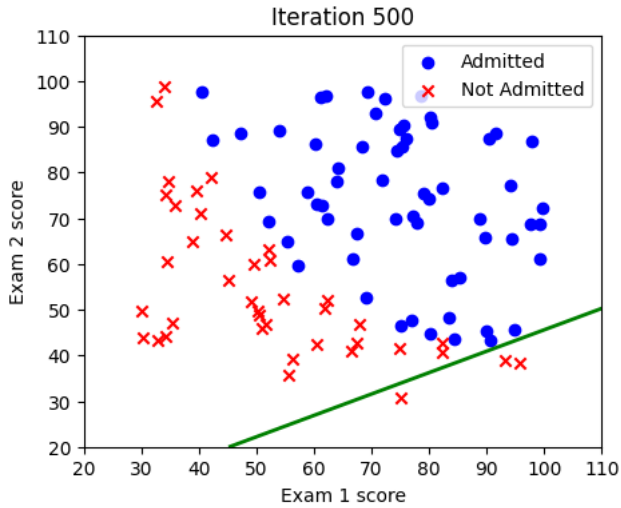





[Iteration 400] cost = 4.794753361216823



[Iteration 500] cost = 7.844443397481126




[Iteration 600] cost = 12.4747540178301




[Iteration 700] cost = 4.779485610522895






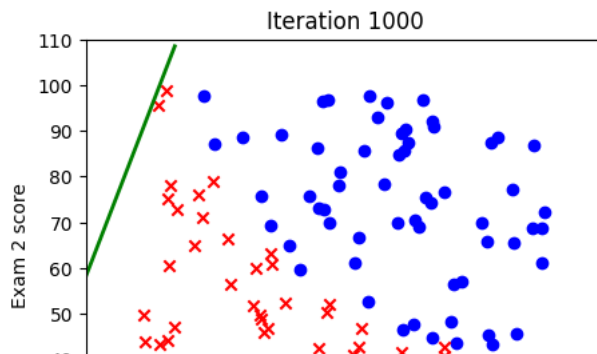
[Iteration 800] cost = 7.821471464743001




[Iteration 900] cost = 12.463702013684587



[Iteration 1000] cost = 4.764248583140833





Purpose:

One-Class SVM (Support Vector Machine) is an **unsupervised learning algorithm** used for **anomaly detection** or **outlier detection**.

Idea: Training Accuracy: 60.0%

It tries to **learn the boundary of a single class** of "normal" data and identifies points that lie **outside this boundary** as anomalies.

How it Works:

1. Maps input data into a **high-dimensional feature space** using a **kernel function** (commonly RBF).
2. Finds a **hyperplane** (or hypersurface) that best **encloses the normal data**.
3. Points outside this boundary are considered **outliers**.

Key Parameters:

- `nu`: An upper bound on the fraction of outliers (controls sensitivity).
- `kernel`: The function used to map data to higher dimensions.

✓ Credit Card Fraud Detection: 3D PCA Visualization

Dataset

- Kaggle's **Credit Card Fraud dataset** with transactions labeled as `0 = normal` and `1 = fraud`.
- Features include `Time`, `Amount`, and 28 anonymized PCA components (`V1` - `V28`).

Steps in the Code

1. **Resampling:** Created a smaller dataset of 2000 transactions with a **70:30 ratio** of normal to fraud for better visualization.
2. **Feature Preparation:** Separated features (`X`) and labels (`y_true`) and **standardized** features using `StandardScaler`.
3. **Dimensionality Reduction:** Applied **PCA** to reduce features to 3 components for 3D plotting.
4. **Visualization:** Plotted normal (blue) and fraud (red) transactions in 3D PCA space to observe patterns and class separation.

Purpose

- Helps visualize the distribution of normal vs. fraudulent transactions and provides intuition before applying anomaly detection or classification models like **One-Class SVM**.

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.preprocessing import StandardScaler
from sklearn.decomposition import PCA
from mpl_toolkits.mplot3d import Axes3D # needed for 3D plots
from sklearn.utils import resample
import kagglehub
from kagglehub import KaggleDatasetAdapter
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import train_test_split
from sklearn.metrics import classification_report, confusion_matrix
import numpy as np
import pandas as pd

# 1. Load dataset from Kaggle
df = kagglehub.dataset_load(
    KaggleDatasetAdapter.PANDAS,
    "mlg-ulb/creditcardfraud",
    "creditcard.csv",
)

# --- 1. Resample to 70:30 ratio ---
fraud = df[df["Class"] == 1]
normal = df[df["Class"] == 0]

# target total size
target_total = 2000
target_fraud = int(0.3 * target_total) # 30% fraud
target_normal = int(0.7 * target_total) # 70% normal
```

```

fraud_resampled = resample(fraud, replace=True, n_samples=target_fraud, random_state=42)
normal_resampled = resample(normal, replace=False, n_samples=target_normal, random_state=42)

df_small = pd.concat([fraud_resampled, normal_resampled]).sample(frac=1, random_state=42)

print(f"Reduced dataset size: {len(df_small)}")
print(df_small["Class"].value_counts(normalize=True))

# --- 2. Features and labels ---
X = df_small.drop("Class", axis=1).values
y_true = df_small["Class"].values # 0 = normal, 1 = fraud

# --- 3. Standardize features ---
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)

# --- 4. Reduce to 3D with PCA ---
pca = PCA(n_components=3, random_state=42)
X_3d = pca.fit_transform(X_scaled)

# --- 5. Plot in 3D ---
fig = plt.figure(figsize=(10, 7))
ax = fig.add_subplot(111, projection='3d')

ax.scatter(X_3d[y_true == 0, 0], X_3d[y_true == 0, 1], X_3d[y_true == 0, 2],
           c='blue', alpha=0.5, label="Normal")
ax.scatter(X_3d[y_true == 1, 0], X_3d[y_true == 1, 1], X_3d[y_true == 1, 2],
           c='red', alpha=0.7, label="Fraud")

ax.set_title("3D PCA Projection of Credit Card Transactions (70:30 ratio)")
ax.set_xlabel("PC 1")
ax.set_ylabel("PC 2")
ax.set_zlabel("PC 3")
ax.legend()
plt.show()

```


Using Colab cache for faster access to the 'creditcardfraud' dataset.
Reduced dataset size: 2000

```

Class
0    0.7
1    0.3
Name: proportion, dtype: float64

```

3D PCA Projection of Credit Card Transactions (70:30 ratio)



One-Class SVM: Primal and Dual Form

The goal of a One-Class SVM is to **find a boundary around the normal data (inliers)** in a high-dimensional space. Formally, we want to find a hyperplane that separates most of the data from the origin in feature space $\phi(x)$.

Primal Formulation

$$\min_{w, \rho, \xi} \quad \frac{1}{2} \|w\|^2 - \rho + \frac{1}{\nu n} \sum_{i=1}^n \xi_i$$

subject to:

$$w^\top \phi(x_i) \geq \rho - \xi_i, \quad \xi_i \geq 0$$

Where:

- w = normal vector of the hyperplane
- ρ = offset / threshold
- ξ_i = slack variables (allow some points to lie outside the boundary)
- $\nu \in (0, 1]$ = upper bound on the fraction of outliers

Intuition:

- Minimize $\|w\|^2 \rightarrow$ keep the hyperplane "tight" around the data
- Maximize $\rho \rightarrow$ push the boundary away from the origin
- Slack ξ_i allows some points to be considered outliers

Dual Formulation

Introducing Lagrange multipliers $\mu_i \geq 0$ for the constraints and solving gives the dual problem:

$$\min_{\mu} \quad \frac{1}{2} \sum_{i,j} \mu_i \mu_j K(x_i, x_j)$$

subject to:

$$0 \leq \mu_i \leq \frac{1}{\nu n}, \quad \sum_i \mu_i = 1$$

Where:

- $K(x_i, x_j) = \phi(x_i)^\top \phi(x_j)$ is the **kernel function**
- μ_i = dual coefficients associated with each training point

Intuition:

- The kernel trick allows computation in high-dimensional space without explicitly mapping $\phi(x)$
- Constraints enforce that only a fraction ν of points can be outliers
- $\sum_i \mu_i = 1$ ensures proper scaling of the hyperplane

✓ Programming Assignment: Kernel One-Class SVM

Overview

In this assignment, you will implement parts of a **Kernel-based One-Class SVM**.

The goal is to learn how kernel methods and dual optimization are used to detect anomalies (outliers vs. inliers).

We provide a partially completed class `KernelOneClassSVM` with:

- Pre-written methods for optimization and training.
- Clear function stubs and docstrings.
- Comments showing exactly what you need to implement.

Your task is to **fill in the missing methods**

Class Provided

The `KernelOneClassSVM` class contains the following methods:

Already Implemented (Do NOT modify)

- `__init__` \rightarrow initializes the model with kernel parameters.

- `fit` → fits the model by solving the dual optimization problem.
- `score_samples` → computes anomaly scores using the decision function.

To Implement (Your Task)

The following functions must be completed by you:

1. `_kernel_function(self, X1, X2)`
 - Compute the kernel matrix between two datasets.
 - Use Gaussian RBF:
2. `_objective(self, mu)`
 - Compute the dual objective function:
3. `_objective_gradient(self, mu)`
 - Compute the gradient of the dual objective:
4. `_compute_rho(self)`
 - Compute the threshold ρ using complementary slackness:
 - For support vectors strictly inside bounds ($0 < \mu_i < 1/(v_n)$):
 - If none exist, take the average over all support vectors.
5. `decision_function(self, X)`
 - Compute decision values for test samples:
6. `predict(self, X)`
 - Predict labels based on decision values:

```
import numpy as np
import matplotlib.pyplot as plt
from scipy.optimize import minimize
from sklearn.metrics import pairwise_distances
from sklearn.metrics import recall_score
import warnings
warnings.filterwarnings('ignore')
```

```
class KernelOneClassSVM:
    def __init__(self, kernel='rbf', gamma=0.1, nu=0.1, tol=1e-6):
        self.kernel = kernel
        self.gamma = gamma
        self.nu = nu
        self.tol = tol

        # These will be set during fitting
        self.X_train = None
        self.mu = None
        self.support_vectors = None
        self.support_indices = None
        self.rho = None
        self.K_train = None

    def _kernel_function(self, X1: np.ndarray, X2: np.ndarray) -> np.ndarray:
        #Compute kernel matrix between X1 and X2.

        #Parameters:
        #-----
        #X1 : array-like, shape (n_samples_1, n_features)
        #X2 : array-like, shape (n_samples_2, n_features)

        #Returns:
        #-----
        #K : array, shape (n_samples_1, n_samples_2)
        #   Kernel matrix
        #   Gaussian RBF kernel:  $\exp(-\gamma ||x_i - x_j||^2)$ 
        # ***** ENTER CODE *****
        distances = pairwise_distances(X1, X2, metric='sqeuclidean')
        K = np.exp(-self.gamma * distances)
        return K
        # ***** END CODE *****

    def _objective(self, mu: np.ndarray) -> float:
```

```

# Parameters
# -----
# mu : numpy.ndarray
#     The vector of dual coefficients or optimization variables, with shape `(n_samples,)`.
# Returns
# -----
# float
#     The scalar value of the dual objective function.
# Implementation Details
# -----
#     The function calculates the objective using the formula:
#      $L(\mu) = \frac{1}{2} \mu^T K \mu$ 
# ***** ENTER CODE *****
objective = 0.5 * np.dot(mu.T, np.dot(self.K_train, mu))
return objective
# ***** END CODE *****

def _objective_gradient(self, mu: np.ndarray) -> np.ndarray:
# Parameters
# -----
# mu : numpy.ndarray
#     The vector of dual coefficients or optimization variables, with shape `(n_samples,)`.
# Returns
# -----
# numpy.ndarray
#     The computed gradient vector, with the same shape as `mu`, `(n_samples,)`.
# # Gradient formula:
#  $\nabla_{\mu} (1/2 * \mu^T K \mu) = K \mu$ 
# ***** ENTER CODE *****
gradient = np.dot(self.K_train, mu)
return gradient
# ***** END CODE *****

# No Changes To be Made in the fit function
def fit(self, X):
# Fit the One-Class SVM model.
# Parameters
# -----
# X : numpy.ndarray, shape (n_samples, n_features)
#     The training data to be used for fitting the model.
# Returns
# -----
# OneClassSVM
#     Returns the instance of the class itself, allowing for method chaining.
# Intitalize the data
X = np.array(X)
self.X_train = X
n_samples = X.shape[0]
# Compute kernel matrix
self.K_train = self._kernel_function(X, X)

# Set up optimization problem
# Variables:  $\mu$  (dual coefficients)
mu_init = np.ones(n_samples) / n_samples # Initialize to satisfy equality constraint

# Bounds:  $0 \leq \mu_i \leq 1/(v*n)$ 
bounds = [(0, 1.0 / (self.nu * n_samples)) for _ in range(n_samples)]

# Equality constraint:  $\mu^T \mathbf{1} = 1$ 
constraint = {
    'type': 'eq',
    'fun': lambda mu: np.sum(mu) - 1.0,
    'jac': lambda mu: np.ones(n_samples)
}

# Solve optimization problem
result = minimize(
    fun=self._objective,
    x0=mu_init,
    method='SLSQP',
    jac=self._objective_gradient,
    bounds=bounds,
    constraints=constraint,
    options={'ftol': self.tol, 'disp': False}
)
if not result.success:
    print(f"Optimization warning: {result.message}")

```

```

self.mu = result.x
self.support_indices = np.where(self.mu > self.tol)[0]
self.support_vectors = X[self.support_indices]
mu_support = self.mu[self.support_indices]
# Compute p using complementary slackness
# For support vectors with  $0 < \mu_i < 1/(v*n)$ , we have  $\langle w, x_i \rangle = p$ 
self._compute_rho()

return self

def _compute_rho(self):
    # Compute the threshold p using complementary slackness conditions.
    # This threshold defines the boundary of the decision function. Data points are
    # classified as inliers or outliers based on their position relative to this threshold.
    # The value of rho is computed based on the support vectors found during optimization.
    # - For support vectors strictly inside bounds ( $0 < \mu_i < 1/(v*n)$ ):
    #    $p = \sum \mu_j K(x_j, x_i)$ 
    # - If none satisfy, take the average decision over all support vectors.
    # Steps:
    # 1. Compute the upper bound for  $\mu_i$ , which is  $1 / (v * n)$ .
    # 2. Identify support vectors where  $\mu_i$  lies strictly between (tol, upper_bound - tol).
    # These are the margin support vectors.
    # 3. If such support vectors exist:
    #   - Pick one (e.g., the first).
    #   - Compute p as  $\sum \mu_j K(x_j, x_i)$ .
    # 4. If none exist:
    #   - Fall back to computing p as the average decision value over all support vectors.
    #   set self.rho no need to return anything
    n_samples = len(self.X_train)
    upper_bound = 1.0 / (self.nu * n_samples)

    # Find support vectors that are strictly between bounds
    mask = (self.mu > self.tol) & (self.mu < upper_bound - self.tol)

    # ***** ENTER CODE *****
    if np.any(mask):
        # Use margin support vectors
        margin_sv_idx = np.where(mask)[0][0] # Take the first one
        self.rho = np.sum(self.mu * self.K_train[:, margin_sv_idx])
    else:
        # Fallback: average over all support vectors
        sv_decision_values = np.sum(self.mu[:, np.newaxis] * self.K_train, axis=0)[self.support_indices]
        self.rho = np.mean(sv_decision_values)
    # ***** END CODE *****

def decision_function(self, X: np.ndarray) -> np.ndarray:
    # Parameters:
    # -----
    # X : array-like, shape (n_samples, n_features)
    #     Test samples
    # Returns:
    # -----
    # decision : array, shape (n_samples,)
    #     Decision function values
    # Decision function formula:
    #  $f(x) = \sum \mu_i K(x, x_i) - p$ 
    # - Positive → inlier
    # - Negative → outlier
    if self.mu is None:
        raise ValueError("Model has not been fitted yet.")

    X = np.array(X)

    # ***** ENTER CODE *****
    K_test = self._kernel_function(X, self.X_train)
    decision = np.dot(K_test, self.mu) - self.rho
    return decision
    # ***** END CODE *****

def predict(self, X: np.ndarray) -> np.ndarray:
    # Parameters:
    # -----
    # X : array-like, shape (n_samples, n_features)
    #     Test samples
    # Returns:

```

```

# y_pred : array, shape (n_samples,)
# Predicted labels (1 for inliers, -1 for outliers)
# - Use decision_function(X)
# - Apply threshold 0:
#   decision >= 0 → inlier (1)
#   decision < 0 → outlier (-1)
# Predicted labels (1 for inliers, -1 for outliers) is an array
# ***** ENTER CODE *****
decision = self.decision_function(X)
y_pred = np.where(decision >= 0, 1, -1)
return y_pred
# ***** END CODE *****
def score_samples(self, X):
    """
    Parameters:
    -----
    X : array-like, shape (n_samples, n_features)
    Returns:
    -----
    scores : array, shape (n_samples,)
        Anomaly scores
    """
    return self.decision_function(X)

X_train, X_test, y_train, y_test = train_test_split(
    X_scaled, y_true, test_size=0.3, random_state=42, stratify=y_true
)

```

grader.check("q5")

q5 passed! 🎉

```

X_train_normal = X_train[y_train == 0]
model = KernelOneClassSVM(kernel="rbf", gamma=0.1, nu=0.05)
model.fit(X_train_normal)

y_pred = model.predict(X_test)
# Convert One-Class SVM outputs (-1 = anomaly, +1 = inlier) → fraud = 1, normal = 0
y_pred_binary = np.where(y_pred == -1, 1, 0)

# 9. Evaluation
print("\n=== Evaluation on Test Set ===")
print(confusion_matrix(y_test, y_pred_binary))
report = classification_report(y_test, y_pred_binary, digits=4)
print(report)
recall = recall_score(y_test, y_pred_binary, average='macro') # or 'weighted', 'micro', None
print("Recall:", recall)

```

```

=== Evaluation on Test Set ===
[[361  59]
 [ 28 152]]

```

	precision	recall	f1-score	support
0	0.9280	0.8595	0.8925	420
1	0.7204	0.8444	0.7775	180
accuracy			0.8550	600
macro avg	0.8242	0.8520	0.8350	600
weighted avg	0.8657	0.8550	0.8580	600

Recall: 0.8519841269841271

grader.check("q6")

q6 passed! 🎉

Do not need to change the below function it is used to - demonstrate the use of a custom **Kernel One-Class SVM** for anomaly detection on credit card fraud data, including 3D visualization using PCA and evaluation on a test set.

```
from sklearn.model_selection import train_test_split
def plot_ocsvm_results_3d(X, y_true, model, title="One-Class SVM Results (3D)"):
    """
    Visualize One-Class SVM results in 3D using PCA for dimensionality reduction.
    Decision boundary is approximated via grid sampling.
    """
    # Reduce to 3D
    pca = PCA(n_components=3, random_state=42)
    X_3d = pca.fit_transform(X)

    fig = plt.figure(figsize=(12, 8))
    ax = fig.add_subplot(111, projection='3d')

    # Scatter plot of inliers/outliers
    inliers_mask = y_true == 0
    outliers_mask = y_true == 1

    ax.scatter(X_3d[inliers_mask, 0], X_3d[inliers_mask, 1], X_3d[inliers_mask, 2],
               c="blue", alpha=0.5, label="Normal")
    ax.scatter(X_3d[outliers_mask, 0], X_3d[outliers_mask, 1], X_3d[outliers_mask, 2],
               c="red", alpha=0.7, label="Fraud")

    # Support vectors (projected in 3D)
    if len(model.support_indices) > 0:
        support_proj = pca.transform(model.X_train[model.support_indices])
        ax.scatter(support_proj[:, 0], support_proj[:, 1], support_proj[:, 2],
                   facecolors="none", edgecolors="black", s=80, linewidth=1.5,
                   label="Support Vectors")

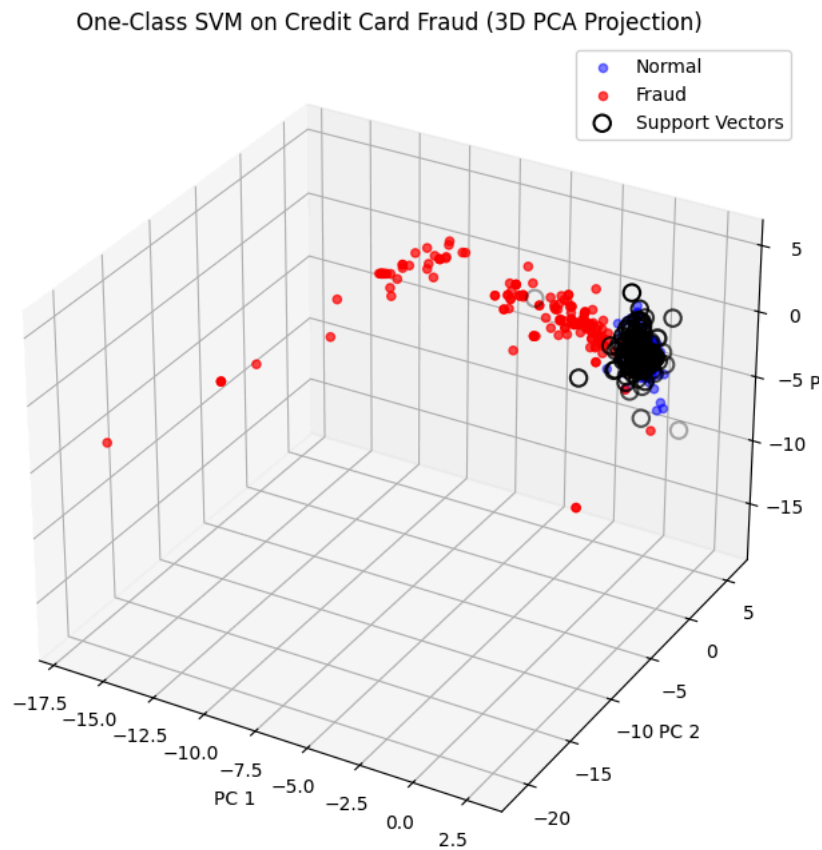
    # Approximate decision boundary using a 3D grid
    grid_size = 20 # reduce if too slow
    x_min, x_max = X_3d[:, 0].min() - 1, X_3d[:, 0].max() + 1
    y_min, y_max = X_3d[:, 1].min() - 1, X_3d[:, 1].max() + 1
    z_min, z_max = X_3d[:, 2].min() - 1, X_3d[:, 2].max() + 1

    xx, yy, zz = np.meshgrid(
        np.linspace(x_min, x_max, grid_size),
        np.linspace(y_min, y_max, grid_size),
        np.linspace(z_min, z_max, grid_size)
    )
    grid_points = np.c_[xx.ravel(), yy.ravel(), zz.ravel()]
    # Project grid back to original space
    grid_original = pca.inverse_transform(grid_points)
    scores = model.decision_function(grid_original)
    scores = scores.reshape(xx.shape)

    # Plot isosurface (decision boundary at 0)
    try:
        from skimage import measure
        verts, faces, _, _ = measure.marching_cubes(scores, level=0)
        verts_transformed = np.c_[xx.ravel()[verts[:, 0].astype(int)],
                                   yy.ravel()[verts[:, 1].astype(int)],
                                   zz.ravel()[verts[:, 2].astype(int)]]
        ax.plot_trisurf(verts_transformed[:, 0], verts_transformed[:, 1],
                        faces, verts_transformed[:, 2], color="cyan", alpha=0.15)
    except Exception as e:
        print("⚠️ Could not render 3D boundary surface:", e)

    ax.set_title(title)
    ax.set_xlabel("PC 1")
    ax.set_ylabel("PC 2")
    ax.set_zlabel("PC 3")
    ax.legend()
    plt.show()

plot_ocsvm_results_3d(X_test, y_test, model,
title="One-Class SVM on Credit Card Fraud (3D PCA Projection)")
```



✓ Pima Indians Diabetes Dataset

The **Pima Indians Diabetes dataset** is a well-known medical dataset from the **National Institute of Diabetes**. It is frequently used for **binary classification tasks** in machine learning.

Dataset Overview

- **Samples (rows):** 768 female patients (age ≥ 21)
- **Features (columns):** 8 health-related attributes
- **Target (label):** Binary outcome
 - 0 → No diabetes
 - 1 → Diabetes

Features

Feature	Description
Pregnancies	Number of times pregnant
Glucose	Plasma glucose concentration
BloodPressure	Diastolic blood pressure (mm Hg)
SkinThickness	Triceps skinfold thickness (mm)
Insulin	2-Hour serum insulin (μ U/ml)
BMI	Body mass index ($\text{weight}/\text{height}^2$)
DiabetesPedigreeFunction	Diabetes pedigree function
Age	Age in years

Data Loading Steps

1. Dataset is **downloaded directly** from GitHub.
2. **Column names** are assigned (since the raw file has none).
3. Data is loaded into a **pandas DataFrame**.

4. **Basic info** is printed:

- Shape of dataset
- Feature names
- Class distribution

5. Data is split into:

- **Features (X)** → NumPy array
- **Target labels (y)** → NumPy array

Summary

- Dataset contains **768 samples** with **8 input features**.
- Target is **binary (0/1)** for diabetes diagnosis.
- No preprocessing applied yet (raw form).

```
import numpy as np
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score, classification_report
from sklearn.preprocessing import LabelEncoder
import requests
from io import StringIO
from typing import Dict, Any, Union

def load_pima_diabetes_data():
    """Load the Pima Indians Diabetes dataset as-is"""
    # UCI dataset URL
    url = "https://raw.githubusercontent.com/jbrownlee/Datasets/master/pima-indians-diabetes.data.csv"

    # Column names for Pima Indians Diabetes dataset
    columns = [
        'pregnancies', 'glucose', 'blood_pressure', 'skin_thickness',
        'insulin', 'bmi', 'diabetes_pedigree', 'age', 'outcome'
    ]

    # Download with requests (disable SSL verify to avoid cert issues)
    print("Downloading Pima Indians Diabetes dataset...")
    response = requests.get(url, verify=False)
    if response.status_code != 200:
        raise Exception(f"Failed to download dataset, status code {response.status_code}")

    # Convert response content to DataFrame
    data = pd.read_csv(StringIO(response.text), names=columns)
    print(f"Successfully loaded {len(data)} samples")

    # Display basic info about the dataset
    print(f"\nDataset shape: {data.shape}")
    print(f"Feature columns: {columns[:-1]}")
    print(f"Target column: {columns[-1]}")
    print(f"Class distribution: {data['outcome'].value_counts().sort_index().values}")

    # Separate features and target - NO preprocessing, use raw data
    X = data.drop('outcome', axis=1).values
    y = data['outcome'].values

    return X, y, columns[:-1]
```

✓ Decision Tree from Scratch

In this section, we will implement **two core functions** used in building a Decision Tree:

1. Gini Impurity

- Measures how mixed the classes are in a node.
- A pure node (all samples same class) has **Gini = 0**.

2. Best Split

- For each feature and threshold, split the data into left/right subsets.

- Compute **weighted Gini impurity** of the split.
- Choose the feature and threshold with the **lowest impurity**.

```

from typing import Dict, Any, Union
class DecisionTreeFromScratch:
    def __init__(self, criterion="gini", max_depth=5, min_samples_split=2):
        self.criterion = criterion
        self.max_depth = max_depth
        self.min_samples_split = min_samples_split
        self.tree = None

    def gini_impurity(self, y: np.ndarray) -> float:
        # Compute the Gini impurity for labels y.
        # Parameters
        # -----
        # y : numpy.ndarray, shape (n_samples,)
        #     The array of class labels for which to compute the impurity.

        # Returns
        # -----
        # float
        #     The computed Gini impurity, a value between 0 and 1. A value of 0 indicates
        #     a pure node (all samples belong to the same class), while a higher value
        #     indicates greater impurity.
        # Formula:  $Gini(y) = 1 - \sum p(c)^2$ 
        if len(y) == 0:
            return 0
        _, counts = np.unique(y, return_counts=True)
        # ***** ENTER CODE *****
        probs = counts / len(y)
        gini = 1 - np.sum(probs ** 2)
        return gini
        # ***** END CODE *****

    def _best_split(self, X: np.ndarray, y: np.ndarray) -> tuple[int, float | None]:
        # Find the best split of data based on Gini impurity.
        # Parameters
        # -----
        # X : numpy.ndarray, shape (n_samples, n_features)
        #     The feature matrix of the dataset.
        # y : numpy.ndarray, shape (n_samples,)
        #     The target labels corresponding to the samples in `X`.

        # Returns
        # -----
        # tuple[int, float | None]
        #     A tuple containing the index of the best feature to split on and the
        #     corresponding threshold value.
        # You have to implement weighted gini impurity
        n_samples, n_features = X.shape
        best_score = float("inf")
        best_feat, best_thresh = None, None

        for feat in range(n_features):
            thresholds = np.unique(X[:, feat])
            for thresh in thresholds:
                left_mask = X[:, feat] <= thresh
                right_mask = ~left_mask

                y_left, y_right = y[left_mask], y[right_mask]

                if len(y_left) == 0 or len(y_right) == 0:
                    continue

                # Weighted Gini impurity
                # ***** ENTER CODE *****
                n_left, n_right = len(y_left), len(y_right)
                gini_left = self.gini_impurity(y_left)
                gini_right = self.gini_impurity(y_right)
                score = (n_left/n_samples) * gini_left + (n_right/n_samples) * gini_right
                # ***** END CODE *****

                if score < best_score:
                    best_score, best_feat, best_thresh = score, feat, thresh

        return best_feat, best_thresh

```

```

def _build(self, X: np.ndarray, y: np.ndarray, depth: int, parent_label: int) -> Dict[str, Any]:
    """Recursively build the decision tree."""
    if len(y) == 0:
        return {"leaf": True, "label": parent_label}

    if (depth >= self.max_depth or len(y) < self.min_samples_split or
        len(np.unique(y)) == 1):
        label = np.bincount(y).argmax()
        return {"leaf": True, "label": label}

    feat, thresh = self._best_split(X, y)
    if feat is None:
        label = np.bincount(y).argmax()
        return {"leaf": True, "label": label}

    left_mask = X[:, feat] <= thresh
    right_mask = ~left_mask
    label = np.bincount(y).argmax()

    return {
        "leaf": False,
        "feature": feat,
        "thresh": thresh,
        "left": self._build(X[left_mask], y[left_mask], depth + 1, label),
        "right": self._build(X[right_mask], y[right_mask], depth + 1, label)
    }

def fit(self, X: np.ndarray, y: np.ndarray) -> None:
    """Fit the decision tree to data X and labels y."""
    parent_label = np.bincount(y).argmax()
    self.tree = self._build(X, y, 0, parent_label)

def _predict_one(self, node: Dict[str, Any], x: np.ndarray) -> int:
    """Predict the label for a single input sample."""
    if node["leaf"]:
        return node["label"]

    if x[node["feature"]] <= node["thresh"]:
        return self._predict_one(node["left"], x)
    else:
        return self._predict_one(node["right"], x)

def predict(self, X: np.ndarray) -> np.ndarray:
    """Predict labels for all samples in dataset X."""
    return np.array([self._predict_one(self.tree, x) for x in X])

```

grader.check("q7")

q7 passed! 🎉

✓ Implementing K-Nearest Neighbors (KNN) from Scratch

In this assignment, you will implement the **K-Nearest Neighbors (KNN)** classification

Background

KNN is a **lazy learning algorithm** that makes predictions for a new data point by looking at the k closest training examples (neighbors) in feature space. The predicted label is decided by **majority vote** among the neighbors.

Steps:

1. Store the training data.
2. For each test sample:
 - Compute the distance to each training sample.
 - Identify the k nearest neighbors.
 - Collect their labels and perform majority voting.

3. Return the predicted label.

Your Task

We provide you with a class `KNNFromScratch` containing function stubs.

You must complete the missing implementation of the following function.

- **predict**: For each test point, compute distances, pick `k` nearest neighbors, and perform majority voting.

```
class KNNFromScratch:
    def __init__(self, k=3, metric="euclidean"):
        self.k = k
        self.metric = metric

    def fit(self, X_train, y_train):
        self.X_train = X_train
        self.y_train = y_train

    def _distance(self, a, b):
        if self.metric == "euclidean":
            return np.sqrt(np.sum((a - b) ** 2))
        elif self.metric == "manhattan":
            return np.sum(np.abs(a - b))
        else:
            raise ValueError(f"Unknown metric: {self.metric}")

    def predict(self, X_test):
        """
        Hints:
        1. For each test sample, compute the distance to every training sample.
        2. Sort the distances and pick the indices of the k nearest neighbors.
        3. Collect the labels of these k nearest neighbors.
        4. Use majority voting to decide the final predicted label.
        5. Store the prediction and repeat for all test samples.
        """
        predictions = []
        for test_point in X_test:
            # ***** ENTER CODE *****
            # Compute distances to all training points
            distances = []
            for train_point in self.X_train:
                dist = self._distance(test_point, train_point)
                distances.append(dist)

            # Get indices of k nearest neighbors
            distances = np.array(distances)
            k_indices = np.argsort(distances)[:self.k]

            # Get labels of k nearest neighbors
            k_labels = self.y_train[k_indices]

            # Majority voting
            unique_labels, counts = np.unique(k_labels, return_counts=True)
            prediction = unique_labels[np.argmax(counts)]
            predictions.append(prediction)
            # ***** END CODE *****
        return np.array(predictions)
```

grader.check("q8")

q8 passed! 🎉

The programming section of this assignment is now complete. From this point onward, you only need to answer the theory questions in the designated Markdown cells.

✓ Do not modify any of the provided code cells beyond this point.
